

CSCI 339: Distributed Systems — Project 1 (Web Server)

Daniel Takeshi Seita, Ziang (Lucky) Zhang
Williams College
{dts1,zz2}.@williams.edu

1 Introduction

A web server listens to connections on a socket to which a client connects. The client can then retrieve files from the server. For instance, a client can run a web server and use it with Firefox to open up the Williams Computer Science homepage. Alternatively, the client can use `telnet` along with a simple text-based protocol to communicate with the server, which will ideally respond with human-readable messages.

The goal of Project 1 was to build a web server allowing these actions. Our server takes two command line arguments: a document root directory and a port number. The document root is the directory containing the files that the client can retrieve. Our web server only supports the HTTP GET request, as that is sufficient to retrieve the contents of a homepage to our server provided that the appropriate files are present in the server's specified document root directory.

Our server was structured using a *multi-threaded approach* and therefore creates a new thread for each incoming connection. When a GET request is sent with the HTTP/1.1 option, after returning results of a request, the server leaves the connection open for some amount of time proportional to the reciprocal of the number of active threads, so more threads result in a shorter timer.

We also strove to make our server user-friendly. To achieve that, we put some special effort in formatting the server's output. Details and examples are in in Section 4.

2 Architectural Overview

Once our server is running, it continually awaits for connections from clients. When a client connects, the server accepts and spawns a new thread for dealing with that specific socket while allowing the original thread to listen for additional connections. The server is able to handle over 1000 connections. For a given accepted connection, our server will perform the following steps once the client has sent in a request. Our server detects this by testing for two carriage returns.

1. Our server parses the request to check for correctness. Currently, only GET requests are supported. If the request is incorrectly formatted according to our rules (explained later), we return a `400 Bad Request`.
2. If the client made an appropriately formatted GET request, we extract the requested file name and search for it. If it is not found, or if the client does not have the correct permissions, we return `404 Not Found` and `403 Forbidden`, respectively.
3. If the file can be found and read, our server sends a response back to the client that includes the contents of the file, as well as HTTP headers and status codes.
4. After responding to the client, the server closes the connection if it was using HTTP/1.0, but keeps the connection open for a bit longer if using HTTP/1.1.

Each of these steps above is described in more detail in the following subsections.

2.1 Parsing the Request

Our server checks the initial request line from the client to see if it starts out with: `GET [directory] HTTP/{1.0, 1.1}`.¹ It does so by extracting the line, splitting it based on whitespace (multiple whitespace are allowed), and checks the first and third tokens. If any of its checks fail, it sends a `400 Bad Request`. Otherwise, it extracts the file name from the second token, which is the `[directory]`, and calls a procedure to validate the file. Additional text after the three tokens is allowed because clients could decide to send more information.

2.2 Checking the File

With a given file path, the server first checks to make sure no instance of `../` is found in the path, since this would mean the client might be accessing parent (or other) files outside the server's document root. The server outputs a `404` error if this happens; otherwise, it moves on to checking the file status. We first use the `stat` function with the file path as input to check for existence, and then used the `S_IROTH` flag to check for file permissions. If either of these checks fail, the server outputs the appropriate error message, which is `404` or `403`, respectively. If the file is deemed to exist and is readable, then the server moves on to sending a response.

2.3 The Server's Response

The server's response includes the following elements: the HTTP status codes, the date, the content type, the content length, the server name, and finally, the requested file itself. At this point, the server responds with an `HTTP/{1.0, 1.1} 200 OK` message, since it did not output any earlier status codes that would have indicated failure. (Our server implements the `200`, `400`, `403`, `404`, and `500` status codes.) It then prints the `Date` header, which displays the current date in

¹The brackets at the end mean either `1.0` or `1.1` can be appended.

the following form: `[DayOfWeek] [Month] [Day] [Hours:Mins:Secs] [Year] ([UTC Time])`. The UTC Time is included so that no one has to try and guess the time zone.

Following that, the server prints out the `Content-Type` header. It detects the type of the file by extracting the file extension from the file name. Currently supported file types are `html/txt`, `jpg`, and `gif`. The server then prints out `Content-Length`, which uses `fseek` to determine the number of bytes of data in the file, and then `Server`, which is the highly creative name `Daniel and Lucky's Server`.

Finally, the actual file is printed!

2.4 HTTP/1.0 and HTTP/1.1

If the client's GET request includes `HTTP/1.0`, then the server simply closes the connection after sending its response. Things get more interesting if using `HTTP/1.1`, because now the server will keep the connection open longer, avoiding the overhead of commencing and closing TCP connection. It keeps the connection open for a thread until its timer goes to zero. Our timeout heuristic makes the server close faster if it is servicing more requests. The formula for computing a given timer is to divide a constant variable by the number of active threads, i.e. $constTime / (\# threads)$. (In our experiments, we set the constant to be 300 seconds.) If a given connection has received a request from the client, the server responds and then resets the timer back to $constTime / (\# threads)$, though the number of threads may have changed while this was happening, so the timer is not guaranteed to be the same as it was before the request. We use a global variable to count the number of active threads and protect it with a lock to prevent race conditions.

2.5 Addendum: Using Threads

We designed our server to create multiple threads for each connection. Of the three major design options available to us (multi-threaded, multi-processed, and event-driven), we decided that using threads was the simplest and cleanest way to achieve our goals, especially since we were building and testing the server using basic GET requests from a

```

-> make && ./server -document_root ~/cs339/csci339_web_server -port 8101
g++ -o server server.cc -lpthread -std=c++0x
Server initializing...
Binded at socket 3.
Started listening.

```

Figure 1: Compiling and running a server.

web browser. We did not see any need to create extra, bulky processors, and vastly preferred low communication context switching and overhead. The risk is with shared resources, but the only item we needed to protect was the count of the number of threads active.

3 Some Problems Encountered

There were several challenging aspects of this project. One was being able to determine how to even get started. Fortunately, the socket programming tutorial from the University of Massachusetts that was provided on the CSCI 339 project website got us up and running with a single-socket server that successfully sent files.

We also did not initially implement re-usability of ports (and sockets). This caused a problem because it created situations when we could successfully run our server on port, say, 8101, and then after closing the server and immediately re-starting it using the exact same command, it would not start! Fortunately, Jeannie directed us to the `setsockopt` function, which allowed us to reuse ports. Her emails also led us to the right direction in timing out our sockets, which were among the last features we implemented.

4 Examples and Testing

In this section, we go over how to use our server and some sample client-server responses. Figure 1 shows the process of compiling the server with `make`, along with a specified document directory. One can immediately see the start of our user-friendly server messages, which clearly explains to the user that the server is binded and ready to accept connections.

Now consider when the user connects to the server via `telnet` and the appropriate port number. Upon successful connection, the user can type

```

-> telnet localhost 8101
Trying 127.0.0.1...
Connected to localhost.
Escape character is '^]'.
GET /index2.html HTTP/1.0

HTTP/1.0 200 OK
Date: Thu Feb 27 22:17:52 2014 (UTC time: 3:17:52)
Content-Type: text/html
Content-Length: 16
Server: Daniel's and Lucky's Server
THIS IS DANIEL!
Connection closed by foreign host.
->

```

Figure 2: What the user sees while connecting to our server.

```

[Socket 5] Opened for new connection.
Number of open connection: 1
[Socket 4] Opened for new connection.
Number of open connection: 2
[Socket 4] Incoming request - GET /index2.html HTTP/1.1
[Socket 4] Message - HTTP/1.1 200 OK
[Socket 4] Message - Date: Thu Feb 27 22:19:14 2014 (UTC time: 3:19:14)
[Socket 4] Message - Content-Type: text/html
[Socket 4] Message - Content-Length: 16
[Socket 4] Message - Server: Daniel's and Lucky's Server
[Socket 4] Successfully sent file /home/cs-students/14dts1/cs339/csci339_web_server/index2.html
[Socket 4] Incoming request -
[Socket 4] Message - HTTP/1.0 400: Bad Request
[Socket 5] Incoming request -
[Socket 5] Message - HTTP/1.0 400: Bad Request
[Socket 6] Opened for new connection.
Number of open connection: 1
[Socket 6] Incoming request -
[Socket 6] Message - HTTP/1.0 400: Bad Request
[Socket 4] Opened for new connection.
Number of open connection: 1
[Socket 4] Incoming request - GET /index2.html HTTP/1.1
[Socket 4] Message - HTTP/1.1 200 OK
[Socket 4] Message - Date: Thu Feb 27 22:23:10 2014 (UTC time: 3:23:10)
[Socket 4] Message - Content-Type: text/html
[Socket 4] Message - Content-Length: 16
[Socket 4] Message - Server: Daniel's and Lucky's Server
[Socket 4] Successfully sent file /home/cs-students/14dts1/cs339/csci339_web_server/index2.html
[Socket 5] Opened for new connection.
Number of open connection: 1
[Socket 4] Opened for new connection.
Number of open connection: 2
[Socket 6] Opened for new connection.
Number of open connection: 3
[Socket 6] Incoming request - no
[Socket 6] Message - HTTP/1.0 400: Bad Request

```

Figure 3: Examples of output from our server, including some examples of bad requests.

in a GET request for a file. This process is shown in Figure 2, with the requested file as `index2.html`. This is a simple file containing just one line: `THIS IS DANIEL!`. The text of the file is shown in the server's response after the appropriate headers: the HTTP status code, the date, the content type, the content length, and the server's name. Notice here that the connection closes immediately after the server's response. Had the request been for HTTP 1.1, the connection would remain until the timer ran out.

Naturally, if we have bad requests, our server should output error messages. While all this is happening, our server is printing out user-friendly messages. A partial segment is shown in Figure 3

5 Questions

Question 1 *Although it wasn't required, how would you go about supporting .htaccess in your server?*

These .htaccesses are text files used to control how a web server can access files in a directory. If our web server had supported these files, then for each file request, it would first extract the IP address of the user. Then, it needs to check the paths leading up to the requested file for .htaccess files to make sure they exist and that they allow clients with that IP address to access it. To change our code, we would parse the file path name by directory (i.e., by the “/” character) and then search sequentially in each directory up to the one containing the actual file itself. Naturally, we return an error and terminate the connection at any point in this process.

Question 2 *Can you think of a scenario in which HTTP/1.0 may perform better than HTTP/1.1? Can you think of a scenario when HTTP/1.1 outperforms HTTP/1.0?*

HTTP/1.0 generally works well when the benefits of immediate parallelism outperform the cost of the slow start and establishing/destroying the TCP connection. One scenario when it would work well would be when the server must send some enormous files and is experiencing significant latency and low bandwidth. In that case, using a single connection per session might not be wise because each one of the gigantic files must be (slowly) processed in order, and the startup and connection costs are lower compared to the cost of the entire session. Thus, HTTP/1.0 is probably better in this case.

What if the server were instead dealing with a large amount of small files, with high bandwidth and low latency. In this hypothetical case, it might make more sense to use HTTP/1.1 to avoid the comparatively more expensive startup and TCP connection costs required by a purely multi-threaded approach that makes a new connection per object. The high bandwidth means each object is processed quickly, which keeps the pipeline moving.

One certainly needs to consider the tradeoffs if the situation is “in between” these extreme cases,

but the above examples should hopefully capture some obvious reasons to use HTTP/1.0 or HTTP/1.1.

6 Conclusion

This was an interesting first project. We started out slow as we spent some time reading over the (very helpful) references given on the assignments page. As mentioned in Section 3, the tutorial on socket programming got us started. Following that, we implemented some of the more minor details, such as generating appropriate HTTP responses, before moving on to managing multiple connections/threads and creating a timeout for the HTTP/1.1 option. The end result is a server that we hope is completely functional, allows multiple clients to connect, correctly implements a timeout to reflect HTTP/1.0 and HTTP/1.1 differences, is user-friendly, can work on multiple browsers (we tested on Firefox and Google Chrome) and is free of unknown bugs.